



# EigenLayer Contest Findings & Analysis Report

2023-07-06

## Table of contents

- [Overview](#)
  - [About C4](#)
  - [Wardens](#)
- [Summary](#)
- [Scope](#)
- [Severity Criteria](#)
- [High Risk Findings \(2\)](#)
  - [\[H-01\] Slot and block number proofs not required for verification of withdrawal \(multiple withdrawals possible\)](#)
  - [\[H-02\] It is impossible to slash queued withdrawals that contain a malicious strategy due to a misplacement of the ++i increment](#)
- [Medium Risk Findings \(2\)](#)
  - [\[M-01\] A staker with verified over-commitment can potentially bypass slashing completely](#)
  - [\[M-02\] A malicious strategy can permanently DoS all currently pending withdrawals that contain it](#)
- [Low Risk and Non-Critical Issues](#)

- [\[L-01\] `computePhase0Eth1DataRoot` always returns an incorrect Merkle tree](#)
- [\[L-02\] `processInclusionProofKeccak` does not work as expected](#)
- [\[L-03\] `merkleizeSha256` doesn't work as expected](#)
- [\[L-04\] `claimableUserDelayedWithdrawals` sometimes returns unclaimable `DelayedWithdrawals` ,so users will see incorrect data](#)
- [\[L-05\] The condition for full withdrawals in the code is different from that in the documentation](#)
- [\[L-06\] Missing validation to a threshold value on full withdrawal](#)
- [\[L-07\] User can stake twice on beacon chain from same eipod, thus losing funds due to same withdrawal credentials](#)
- [Gas Optimizations](#)
  - [G-01 Optimize `merkleizeSha256` function for gas-efficiency](#)
  - [G-02 Use unchecked arithmetic in `processInclusionProofKeccak` and `processInclusionProofSha256` functions](#)
- [Disclosures](#)



## Overview



## About C4

Code4rena (C4) is an open organization consisting of security researchers, auditors, developers, and individuals with domain expertise in smart contracts.

A C4 audit is an event in which community participants, referred to as Wardens, review, audit, or analyze smart contract logic in exchange for a bounty provided by sponsoring projects.

During the audit outlined in this document, C4 conducted an analysis of the EigenLayer smart contract system written in Solidity. The audit took place between April 27 - May 4 2023.



## Wardens

48 Wardens contributed reports to the EigenLayer Audit:

1. [OxSmartContract](#)
2. [OxTheCOrder](#)
3. OxWaitress
4. [Oxnev](#)
5. [8olidity](#)
6. [ABA](#)
7. [Aymen0909](#)
8. [CoOnan](#)
9. [Cyfrin](#) ([PatrickAlphaC](#), [giovannidisiena](#), [hansfrieze](#), [OKage](#), alexroan and [carlitox477](#))
10. Dug
11. HaipIs
12. Josiah
13. [MiloTruck](#)
14. [QiuhaoLi](#)
15. RaymondFam
16. ReyAdmirado
17. [Ruhum](#)
18. SpicyMeatball
19. [ToonVH](#)
20. [bin2chen](#)
21. btk
22. bughunter007
23. [bytes032](#)
24. clayj
25. d3e4
26. [evmboi32](#)
27. [ihtishamsudo](#)

- 28. [itsmeSTYJ](#)
- 29. jasonxiale
- 30. [juancito](#)
- 31. libratus
- 32. [naman1778](#)
- 33. neutiyoo
- 34. niser93
- 35. pontifex
- 36. rvierdiiev
- 37. said
- 38. sashik\_eth
- 39. tonisives
- 40. [turvy\\_fuzz](#)
- 41. [volodya](#)
- 42. [windowhan001](#)
- 43. yjrwkk

This audit was judged by [Alex the Entrepreneurd](#).

Final report assembled by thebrittfactor.



## Summary

The C4 analysis yielded an aggregated total of 4 unique vulnerabilities. Of these vulnerabilities, 2 received a risk rating in the category of HIGH severity and 2 received a risk rating in the category of MEDIUM severity.

Additionally, C4 analysis included 15 reports detailing issues with a risk rating of LOW severity or non-critical. There were also 13 reports recommending gas optimizations.

All of the issues presented here are linked back to their original finding.



# Scope

The code under review can be found within the [C4 EigenLayer Audit repository](#), and is composed of 24 smart contracts written in the Solidity programming language and includes 1393 lines of Solidity code.



## Severity Criteria

C4 assesses the severity of disclosed vulnerabilities based on three primary risk categories: high, medium, and low/non-critical.

High-level considerations for vulnerabilities span the following key areas when conducting assessments:

- Malicious Input Handling
- Escalation of privileges
- Arithmetic
- Gas use

For more information regarding the severity criteria referenced throughout the submission review process, please refer to the documentation provided on [the C4 website](#), specifically our section on [Severity Categorization](#).



## High Risk Findings (2)



[H-01] Slot and block number proofs not required for verification of withdrawal (multiple withdrawals possible)

*Submitted by [OxTheC0der](#), also found by [windowhan001](#) and [volodya](#)*

Since this is a vulnerability that involves multiple in-scope contracts and leads to more than one impact, let's start with a bug description from bottom to top.



**Library** Merkle

The methods [verifyInclusionSha256\(proof, root, leaf, index\)](#) and [verifyInclusionKeccak\(proof, root, leaf, index\)](#) will always return `true` if `proof.length < 32` (e.g. empty proof) and `leaf == root`. Although this might

be intended behaviour, I see no use case for empty proofs and would require non-empty proofs at the library level. As of now, the user of the library is responsible to enforce non-zero proofs.



## Library BeaconChainProofs

The method [verifyWithdrawalProofs\(beaconStateRoot, proofs, withdrawalFields\)](#), which relies on multiple calls to [Merkle.verifyInclusionSha256\(proof, root, leaf, index\)](#), does not require a minimum length of `proofs.slotProof` and `proofs.blockNumberProof`. As a consequence, considering a valid set of `(beaconStateRoot, proofs, withdrawalFields)`, the method will still succeed with **empty** slot and block number proofs, i.e. the `proofs` can be modified in the following way:

```
proofs.slotProof = bytes(""); // empty slot proof
proofs.slotRoot = proofs.blockHeaderRoot; // make leaf == root

proofs.blockNumberProof = bytes(""); // empty k
proofs.blockNumberRoot = proofs.executionPayloadRoot; // make leaf == root
```

As a consequence, we can take a perfectly valid withdrawal proof and re-create the proof for the same withdrawal with a **different** slot and block number (according to the code above) that will still be accepted by the [verifyWithdrawalProofs\(beaconStateRoot, proofs, withdrawalFields\)](#) method.



## Contract EigenPod

The method [verifyAndProcessWithdrawal\(withdrawalProofs, ...\)](#), which relies on a call to [BeaconChainProofs.verifyWithdrawalProofs\(beaconStateRoot, proofs, withdrawalFields\)](#), is impacted by a modified - but still valid - withdrawal proof in two ways.

First, the modifier

[proofIsForValidBlockNumber\(Endian.fromLittleEndianUint64\(withdrawalProofs.blockNumberRoot\)\)](#) makes sure that the **block number** being proven is greater/newer than the `mostRecentWithdrawalBlockNumber`. In our case, `blockNumberRoot = executionPayloadRoot` and depending on the actual value of `executionPayloadRoot`, the `proofIsForValidBlockNumber` can be bypassed as

shown in the test, see any PoC test case. As a consequence, old withdrawal proofs could be re-used with an empty `blockNumberProof` to withdraw the same funds more than once.

Second, the sub-method [`\_processPartialWithdrawal\(withdrawalHappenedSlot, ...\)`](#) requires that a `slot` is only used once. In our case, `slotRoot = blockHeaderRoot` leads to a [different slot](#) than suggested by the original proof. Therefore, a withdrawal proof can be re-used with an empty `slotProof` to do the same partial withdrawal twice, see PoC. Depending on the actual value of `blockHeaderRoot`, a full withdrawal, instead of a partial withdrawal, will be done according to the [condition in L354](#).

Insufficient validation of proofs allows multiple withdrawals, i.e. theft of funds.



## Proof of Concept

The changes to the `EigenPod` test cases below demonstrate the following outcomes:

**testFullWithdrawalProof:**

[`BeaconChainProofs.verifyWithdrawalProofs\(beaconStateRoot, proofs, withdrawalFields\)`](#) still succeeds on empty slot and block number proofs.

**testFullWithdrawalFlow:** [`EigenPod.verifyAndProcessWithdrawal\(withdrawalProofs, ...\)`](#) allows full withdrawal with empty slot and block number proofs.

**testPartialWithdrawalFlow:**

[`EigenPod.verifyAndProcessWithdrawal\(withdrawalProofs, ...\)`](#) allows partial withdrawal with empty slot and block number proofs.

**testProvingMultipleWithdrawalsForSameSlot:**

[`EigenPod.verifyAndProcessWithdrawal\(withdrawalProofs, ...\)`](#) allows partial withdrawal of the same funds twice due to different `slotRoot` in original and modified proof.

The

[`proofsForValidBlockNumber\(Endian.fromLittleEndianUint64\(withdrawalProofs.blockNumberRoot\)\)`](#) modifier is bypassed (see `blockNumberRoot`) in the latter three of the above test cases.

Apply the following *diff* to your `src/test/EigenPod.t.sol` and run the tests with `forge test --match-contract EigenPod`:

```

diff --git a/src/test/EigenPod.t.sol b/src/test/EigenPod.t.sol
index 31e6a58..5242def 100644
--- a/src/test/EigenPod.t.sol
+++ b/src/test/EigenPod.t.sol
@@ -260,7 +260,7 @@ contract EigenPodTests is ProofParsing, EigenPod {

    function testFullWithdrawalProof() public {
        setJSON("./src/test/test-data/fullWithdrawalProof.json"
-       BeaconChainProofs.WithdrawalProofs memory proofs = _get
+       BeaconChainProofs.WithdrawalProofs memory proofs = _get
        withdrawalFields = getWithdrawalFields();
        validatorFields = getValidatorFields();

@@ -281,7 +281,7 @@ contract EigenPodTests is ProofParsing, EigenPod {

        // ./solidityProofGen "WithdrawalFieldsProof" 61336 226
        setJSON("./src/test/test-data/fullWithdrawalProof.json"
-       BeaconChainProofs.WithdrawalProofs memory withdrawalProo
+       BeaconChainProofs.WithdrawalProofs memory withdrawalProo
        bytes memory validatorFieldsProof = abi.encodePacked(get
        withdrawalFields = getWithdrawalFields();
        validatorFields = getValidatorFields();

@@ -317,7 +317,7 @@ contract EigenPodTests is ProofParsing, EigenPod {
        //generate partialWithdrawalProofs.json with:
        // ./solidityProofGen "WithdrawalFieldsProof" 61068 656
        setJSON("./src/test/test-data/partialWithdrawalProof.js
-       BeaconChainProofs.WithdrawalProofs memory withdrawalProo
+       BeaconChainProofs.WithdrawalProofs memory withdrawalProo
        bytes memory validatorFieldsProof = abi.encodePacked(get

        withdrawalFields = getWithdrawalFields();

@@ -346,21 +346,22 @@ contract EigenPodTests is ProofParsing, EigenPod {

        /// @notice verifies that multiple partial withdrawals can
        function testProvingMultipleWithdrawalsForSameSlot(/*uint256
-       IEigenPod newPod = testPartialWithdrawalFlow();
+       IEigenPod newPod = testPartialWithdrawalFlow(); // uses

-       BeaconChainProofs.WithdrawalProofs memory withdrawalProo
+       BeaconChainProofs.WithdrawalProofs memory withdrawalProo
        bytes memory validatorFieldsProof = abi.encodePacked(get
        withdrawalFields = getWithdrawalFields();
        validatorFields = getValidatorFields();

-       cheats.expectRevert(bytes("EigenPod._processPartialWith

```



```

+         // do not expect revert anymore due to different 'slotF
+         //cheats.expectRevert(bytes("EigenPod._processPartialWi
newPod.verifyAndProcessWithdrawal(withdrawalProofs, val
    }

    /// @notice verifies that multiple full withdrawals for a s
function testDoubleFullWithdrawal() public {
-     IEigenPod newPod = testFullWithdrawalFlow();
-     BeaconChainProofs.WithdrawalProofs memory withdrawalPro
+     IEigenPod newPod = testFullWithdrawalFlow(); // uses SP
+     BeaconChainProofs.WithdrawalProofs memory withdrawalPro
bytes memory validatorFieldsProof = abi.encodePacked(get
withdrawalFields = getWithdrawalFields();
    validatorFields = getValidatorFields();
@@ -759,8 +760,11 @@ contract EigenPodTests is ProofParsing, Eig
    return proofs;
}

+     uint256 internal constant FULL_PROOF = 0;
+     uint256 internal constant SKIP_SLOT_BLOCK_PROOF = 1;
+
    /// @notice this function just generates a valid proof so t
-     function _getWithdrawalProof() internal returns(BeaconChair
+     function _getWithdrawalProof(uint256 proofType) internal re
        //make initial deposit
        cheats.startPrank(podOwner);
        eigenPodManager.stake{value: stakeAmount}(pubkey, signa
@@ -773,9 +777,9 @@ contract EigenPodTests is ProofParsing, Eige
        beaconChainOracle.setBeaconChainStateRoot(beaconSta
        bytes32 blockHeaderRoot = getBlockHeaderRoot();
        bytes32 blockBodyRoot = getBlockBodyRoot();
-         bytes32 slotRoot = getSlotRoot();
-         bytes32 blockNumberRoot = getBlockNumberRoot();
+         bytes32 slotRoot = (proofType == FULL_PROOF) ? getS
        bytes32 executionPayloadRoot = getExecutionPayloadF
+         bytes32 blockNumberRoot = (proofType == FULL_PROOF)

@@ -786,9 +790,9 @@ contract EigenPodTests is ProofParsing, Eige
        BeaconChainProofs.WithdrawalProofs memory proofs =
            abi.encodePacked(getBlockHeaderProof()),
            abi.encodePacked(getWithdrawalProof()),
-         abi.encodePacked(getSlotProof()),
+         (proofType == FULL_PROOF) ? abi.encodePacked(get
            abi.encodePacked(getExecutionPayloadProof()),

```

```
-         abi.encodePacked(getBlockNumberProof()),
+         (proofType == FULL_PROOF) ? abi.encodePacked(ge
uint64(blockHeaderRootIndex),
uint64(withdrawalIndex),
blockHeaderRoot,
```

We can see that **all** the test cases are still passing, whereby the following ones are confirming the aforementioned outcomes:

```
[PASS] testFullWithdrawalFlow():(address) (gas: 28517915)
[PASS] testFullWithdrawalProof() (gas: 13185538)
[PASS] testPartialWithdrawalFlow():(address) (gas: 28679149)
[PASS] testProvingMultipleWithdrawalsForSameSlot() (gas: 4550228
```



## Tools Used

VS Code, Foundry



## Recommended Mitigation Steps

Require a minimum length (tree height) for the slot and block number proofs in [BeaconChainProofs.verifyWithdrawalProofs\(beaconStateRoot, proofs, withdrawalFields\)](#).

At least require non-empty proofs according to the following *diff*:

```
diff --git a/src/contracts/libraries/BeaconChainProofs.sol b/src
index b4129bf..119baf2 100644
--- a/src/contracts/libraries/BeaconChainProofs.sol
+++ b/src/contracts/libraries/BeaconChainProofs.sol
@@ -259,6 +259,10 @@ library BeaconChainProofs {
    "BeaconChainProofs.verifyWithdrawalProofs: withdraw
require(proofs.executionPayloadProof.length == 32 * (BE
    "BeaconChainProofs.verifyWithdrawalProofs: executio
+   require(proofs.slotProof.length >= 32,
+   "BeaconChainProofs.verifyWithdrawalProofs: slotProc
+   require(proofs.blockNumberProof.length >= 32,
+   "BeaconChainProofs.verifyWithdrawalProofs: blockNun

/**
```

\* Computes the `block_header_index` relative to the beacon

**Alternative:** Non-empty proofs can also be required in the `Merkle` library.



Assessed type

Invalid Validation

[sorrynotsorry \(lookout\) commented:](#)

Well demonstrated with referrable code snippets, hyperlinks, and coded POC.  
Marking as HQ.

[Sidu28 \(EigenLayer\) confirmed](#)

[Alex the Entrepreneurd \(judge\) commented:](#)

The Warden has shown how, due to a lack of length check, an empty proof could be provided; which would pass validation.

This is an example of how a lack of a check can be chained into a proper exploit, and because the proof will pass, funds can be stolen.

For these reasons I agree with High Severity.



[H-02] It is impossible to slash queued withdrawals that contain a malicious strategy due to a misplacement of the `++i` increment

Submitted by [juancito](#), also found by [yjrwwk](#), [pontifex](#), [evmboi32](#), [bin2chen](#), [sashik\\_eth](#), [Ruhum](#), [MiloTruck](#), [SpicyMeatball](#), and [volodya](#).

`StrategyManager::slashQueuedWithdrawal()` contains an `indicesToSkip` parameter to skip malicious strategies, as documented in the [function definition](#):

so that, e.g., if the slashed `QueuedWithdrawal` contains a malicious strategy in the `strategies` array which always reverts on calls to its 'withdraw' function,

then the malicious strategy can be skipped (with the shares in effect “burned”), while the non-malicious strategies are still called as normal.

The problem is, the function does not work as expected, and `indicesToSkip` is ignored. If the queued withdrawal contains a malicious strategy, it will make the slash always revert.

Owners won't be able to slash queued withdrawals that contain a malicious strategy.

An adversary can take advantage of this and create withdrawal queues that won't be able to be slashed, completely defeating the slash system. The adversary can later complete the withdrawal.



## Proof of Concept

The `++i;` statement in `StrategyManager::slashQueuedWithdrawal()` is misplaced. It is only executed on the `else` statement:

```
// keeps track of the index in the `indicesToSkip` array
uint256 indicesToSkipIndex = 0;

uint256 strategiesLength = queuedWithdrawal.strategies.length
for (uint256 i = 0; i < strategiesLength;) {
    // check if the index i matches one of the indices speci
    if (indicesToSkipIndex < indicesToSkip.length && indices
        unchecked {
            ++indicesToSkipIndex;
        }
    } else {
        if (queuedWithdrawal.strategies[i] == beaconChainETH
            //withdraw the beaconChainETH to the recipie
            _withdrawBeaconChainETH(queuedWithdrawal.deposit
        } else {
            // tell the strategy to send the appropriate amc
            queuedWithdrawal.strategies[i].withdraw(recipier
        }
        unchecked {
            ++i; // @audit
        }
    }
}
```

[Link to code](#)

Let's suppose that the owner tries to slash a queued withdrawal, and wants to skip the first strategy (index 0) because it is malicious and makes the whole transaction revert.

1. It defines `indicesToSkipIndex = 0`.
2. It enters the `for` loop starting at `i = 0`.
3. `if (indicesToSkipIndex < indicesToSkip.length && indicesToSkip[indicesToSkipIndex] == i)` will be true: `0 < 1 && 0 == 0`.
4. It increments `++indicesToSkipIndex;` to "skip" the malicious strategy, so now `indicesToSkipIndex = 1`.
5. It goes back to the `for` loop. But `i` hasn't been modified, so still `i = 0`.
6. `if (indicesToSkipIndex < indicesToSkip.length && indicesToSkip[indicesToSkipIndex] == i)` will be false now: `1 < 1 && 0 == 0`.
7. It will enter the `else` statement and attempt to slash the strategy anyway.
8. If the strategy is malicious, it will revert, making it impossible to slash.
9. The adversary can later complete the withdrawal.



## POC Test

This test shows how the `indicesToSkip` parameter is completely ignored.

For the sake of simplicity of the test, it uses a normal strategy; which will be slashed, proving that it ignores the `indicesToSkip` parameter and it indeed calls

```
queuedWithdrawal.strategies[i].withdraw()
```

A malicious strategy that makes `withdraw()` revert, would be to make the whole transaction revert (not shown on this test but easily checkable as the [function won't catch it](#)).

Add this test to `src/tests/StrategyManagerUnit.t.sol` and run `forge test -m "testSlashQueuedWithdrawal_IgnoresIndicesToSkip"`.

```
function testSlashQueuedWithdrawal_IgnoresIndicesToSkip() external {
    address recipient = address(this);
```

```

uint256 depositAmount = 1e18;
uint256 withdrawalAmount = depositAmount;
bool undelegateIfPossible = false;

// Deposit into strategy and queue a withdrawal
(IStrategyManager.QueuedWithdrawal memory queuedWithdrawal) =
    testQueueWithdrawal_ToSelf_NotBeaconChainETH(depositAmount, withdrawalAmount);

// Slash the delegatedOperator
slasherMock.freezeOperator(queuedWithdrawal.delegatedAddress);

// Keep track of the balance before the slash attempt
uint256 balanceBefore = dummyToken.balanceOf(address(recipient));

// Assert that the strategies array only has one element
assertEq(queuedWithdrawal.strategies.length, 1);

// Set `indicesToSkip` so that it should ignore the only strategy
// As it's the only element, its index is `0`
uint256[] memory indicesToSkip = new uint256[](1);
indicesToSkip[0] = 0;

// Call `slashQueuedWithdrawal()`
// This should not try to slash the only strategy the queue has
// But in fact it ignores `indicesToSkip` and attempts to slash the strategy
cheats.startPrank(strategyManager.owner());
strategyManager.slashQueuedWithdrawal(recipient, queuedWithdrawal);
cheats.stopPrank();

uint256 balanceAfter = dummyToken.balanceOf(address(recipient));

// The `indicesToSkip` was completely ignored, and the balance increased
// It can be asserted due to the fact that it increased by the withdrawal amount
require(balanceAfter == balanceBefore + withdrawalAmount);
}

```



## Recommended Mitigation Steps

Place the `++i` outside of the if/else statement. This way, it will increment each time the loop runs.

```

for (uint256 i = 0; i < strategiesLength; i++) {
    // check if the index i matches one of the indices specified in indicesToSkip
    if (indicesToSkipIndex < indicesToSkip.length && indicesToSkip[indicesToSkipIndex] == i) {
        continue;
    }
}

```

```

        unchecked {
            ++indicesToSkipIndex;
        }
    } else {
        if (queuedWithdrawal.strategies[i] == beaconChainETH
            //withdraw the beaconChainETH to the recipient
            _withdrawBeaconChainETH(queuedWithdrawal.deposit
        } else {
            // tell the strategy to send the appropriate amount
            queuedWithdrawal.strategies[i].withdraw(recipient)
        }
    }
    unchecked {
        ++i;
    }
}

+ unchecked {
+     ++i;
+ }
}

```



## Assessed type

Loop

[sorrynotsorry \(lookout\) commented:](#)

The issue is well demonstrated, properly formatted, and contains a coded POC.  
Marking as HQ.

[Sidu28 \(EigenLayer\) confirmed](#)

[Alex the Entrepreneurd \(judge\) commented:](#)

The Warden has shown how, due to incorrect placement of the loop increment, malicious strategies cannot be skipped when slashing queued withdrawals.

Because this breaks a core functionality of the contracts, which will also cause a loss of funds, I agree with High Severity.

Mitigation is straightforward.



## Medium Risk Findings (2)



[M-01] A staker with verified over-commitment can potentially bypass slashing completely

Submitted by [Cyfrin](#), also found by [Josiah](#), [Qiu hao Li](#), and [Raymond Fam](#)

<https://github.com/code-423n4/2023-04-eigenlayer/blob/5e4872358cd2bda1936c29f460ece2308af4def6/src/contracts/core/StrategyManager.sol#L197>

<https://github.com/code-423n4/2023-04-eigenlayer/blob/5e4872358cd2bda1936c29f460ece2308af4def6/src/contracts/core/StrategyManager.sol#L513>

In EigenLayer, watchers submit over-commitment proof in the event a staker's balance on the Beacon chain falls below the minimum restaked amount per validator. In such a scenario, stakers' shares are decreased by the restaked amount. Note, that when a full withdrawal is processed, stakers' deducted shares are credited back to allow for a planned withdrawal on EigenLayer

If such a staker has delegated to an operator who gets slashed on EigenLayer, there is a possibility that this staker completely bypasses any slashing penalties. If overcommitment reduced the shares in the stakers account to 0, there is nothing available for governance to slash.

It is reasonable to assume that governance calls `StrategyManager::slashShares` to reduce a percentage of shares (penalty) from all stakers who delegated to a slashed operator and then resets the frozen status of operator by calling `ISlasher::resetFrozenStatus`.

By simply unstaking on the beacon chain AFTER slashing is complete (and operator is unfrozen), an over-committed staker can simply unstake on beacon chain and get back their shares that were deducted when over-commitment was recorded (refer to PoC). Note: these shares have not faced any slashing penalties.

An over-committed staker can avoid being slashed in a scenario, in which their stake should be subject to slashing and so we evaluate the severity to **MEDIUM**.





## Proof of Concept

Consider the following scenario with a chain of events in this order:

1. Alice stakes 32 ETH and gets corresponding shares in `BeaconEthStrategy`.
2. After some time, Alice gets slashed on the Beacon Chain and her current balance on the Beacon Chain is now less than what she restaked on EigenLayer.
3. An observer will submit a proof via `EigenPod::verifyOverCommittedStake` and Alice's shares are now decreased to 0 (Note: this will be credited back to Alice when she withdraws from the Beacon Chain using `EigenPod::verifyAndProcessWithdrawal`).
4. Next, Alice's operator gets slashed and her account gets frozen.
5. Governance slashes Alice, along with all stakers who delegated to slashed operator (there is nothing to slash since Alice's shares are currently 0).
6. After slashing everyone, governance resets frozen status of operator by calling `ISlasher::resetFrozenStatus`.
7. Alice now unstakes on the Beacon Chain and gets a credit of shares that were earlier deducted while recording over-commitment.
8. Alice queues a withdrawal request and completes the withdrawal, without facing any slashing penalty.



## Recommended Mitigation Steps

Over-commitment needs to be accounted for when slashing, such that a staker is being slashed; not just shares in their account, but also the amount that is temporarily debited while recording over-commitment.

Consider adding a mapping to `StrategyManager` that keeps track of over-commitment for each staker and take that into consideration while slashing in `StrategyManager::slashShares`.

### [Alex the Entrepreneurd \(judge\) commented:](#)

The Warden has shown how, due to an incorrect assumption, it's possible for an over-committed staker to avoid a slashing.

The finding requires multiple external requirements:

- Being over-committed (perhaps inactivity leak or ETH2 slashing).
- Having delegated to an operator that will be slashed.

For these reasons I agree with Medium Severity.

[Sidu28 \(EigenLayer\) confirmed via duplicate issue 210](#)

🔗

[M-02] A malicious strategy can permanently DoS all currently pending withdrawals that contain it

*Submitted by [ABA](#), also found by [juancito](#), [ToonVH](#), [OxWaitress](#), [ABA](#), [MiloTruck](#), [rvierdiiev](#), [8olidity](#), [bughunter007](#), and [bytes032](#).*

In order to withdraw funds from the project a user has to:

1. queue a withdrawal (via `queueWithdrawal` ).
2. complete a withdrawal (via `completeQueuedWithdrawal(s)` ).

Queuing a withdrawal, via `queueWithdrawal` modifies all internal accounting to reflect this:

<https://github.com/code-423n4/2023-04-eigenlayer/blob/main/src/contracts/core/StrategyManager.sol#L345-L346>

```
// modify delegated shares accordingly, if applicable
delegation.decreaseDelegatedShares(msg.sender, strategie
```

<https://github.com/code-423n4/2023-04-eigenlayer/blob/main/src/contracts/core/StrategyManager.sol#L370>

```
if (_removeShares(msg.sender, strategyIndexes[strategyIndex]
```

and saves the withdrawal hash in order to be used in

```
completeQueuedWithdrawal(s) .
```

<https://github.com/code-423n4/2023-04->

[eigenlayer/blob/main/src/contracts/core/StrategyManager.sol#L400-L415](https://github.com/code-423n4/2023-04-eigenlayer/blob/main/src/contracts/core/StrategyManager.sol#L400-L415)

```
        queuedWithdrawal = QueuedWithdrawal({
            strategies: strategies,
            shares: shares,
            depositor: msg.sender,
            withdrawerAndNonce: withdrawerAndNonce,
            withdrawalStartBlock: uint32(block.number),
            delegatedAddress: delegatedAddress
        });
    }

    // calculate the withdrawal root
    bytes32 withdrawalRoot = calculateWithdrawalRoot(queuedWithdrawal);

    // mark withdrawal as pending
    withdrawalRootPending[withdrawalRoot] = true;
```

In other words, it is final (as there is no `cancelWithdrawal` mechanism implemented). When executing `completeQueuedWithdrawal(s)`, the `withdraw` function of the strategy is called (if `receiveAsTokens` is set to true).

<https://github.com/code-423n4/2023-04->

[eigenlayer/blob/main/src/contracts/core/StrategyManager.sol#L786-L789](https://github.com/code-423n4/2023-04-eigenlayer/blob/main/src/contracts/core/StrategyManager.sol#L786-L789)

```
// tell the strategy to send the appropriate amount of funds
queuedWithdrawal.strategies[i].withdraw(
    msg.sender, tokens[i], queuedWithdrawal.shares[i]
);
```

In this case, a malicious strategy can always revert, blocking the user from retrieving his tokens. If a user sets `receiveAsTokens` to `false`, the other case, then the tokens will be added as shares to the delegation contract

<https://github.com/code-423n4/2023-04->

<eigenlayer/blob/main/src/contracts/core/StrategyManager.sol#L797-L800>

```
for (uint256 i = 0; i < strategiesLength;) {
    _addShares(msg.sender, queuedWithdrawal.strategies[i], c
    unchecked {
        ++i;
    }
}
```

<https://github.com/code-423n4/2023-04->

<eigenlayer/blob/main/src/contracts/core/StrategyManager.sol#L646-L647>

```
// if applicable, increase delegated shares accordingly
delegation.increaseDelegatedShares(depositor, strategy,
```

This still poses a problem because the `increaseDelegatedShares` function's counterpart, `decreaseDelegatedShares`, is only callable by the strategy manager (for users to indirectly get their rewards worth back, via this workaround).

<https://github.com/code-423n4/2023-04->

<eigenlayer/blob/main/src/contracts/core/DelegationManager.sol#L168-L179>

```
/**
 * @notice Decreases the `staker`'s delegated shares in each
 * @dev Callable only by the StrategyManager
 */
function decreaseDelegatedShares(
    address staker,
    IStrategy[] calldata strategies,
    uint256[] calldata shares
)
    external
    onlyStrategyManager
{
```

But in `StrategyManager`, there are only 3 cases where this is called, none of which is beneficial for the user:

- [`recordOvercommittedBeaconChainETH`](#) - slashes user rewards in certain conditions
- [`queueWithdrawal`](#) - accounting side effects have already been done, will fail in `_removeShares`
- [`slashShares`](#) - slashes user rewards in certain conditions

In other words, the workaround (of setting `receiveAsTokens` to `false` in `completeQueuedWithdrawal(s)` ) just leaves the funds accounted and stuck in `DelegationManager`.

In both cases, all shares/tokens associated with other strategies in the pending withdrawal are permanently blocked.



## Proof of Concept

<https://github.com/code-423n4/2023-04-eigenlayer/blob/main/src/contracts/core/StrategyManager.sol#L786-L789>



## Recommended Mitigation Steps

Add a `cancelQueuedWithdrawal` function that will undo the accounting and cancel out any dependency on the malicious strategy. Although this solution may create a front-running opportunity for when their withdrawal will be slashed via `slashQueuedWithdrawal`, there may exist workarounds to this.

Another possibility is to implement a similar mechanism to how `slashQueuedWithdrawal` treats malicious strategies: adding a list of strategy indices to skip.

<https://github.com/code-423n4/2023-04-eigenlayer/blob/main/src/contracts/core/StrategyManager.sol#L532-L535>  
<https://github.com/code-423n4/2023-04-eigenlayer/blob/main/src/contracts/core/StrategyManager.sol#L560-L566>

```
* @param indicesToSkip Optional input parameter -- indices
* so that, e.g., if the slashed QueuedWithdrawal contains a
* then the malicious strategy can be skipped (with the shar
*/
```

...

```
for (uint256 i = 0; i < strategiesLength;) {  
    // check if the index i matches one of the indices s  
    if (indicesToSkipIndex < indicesToSkip.length && inc  
        unchecked {  
        ++indicesToSkipIndex;  
    }  
} else {
```



## Assessed type

DoS

### Sidu28 (EigenLayer) disagreed with severity and commented:

The description of lack of check here is accurate. There is zero actual impact; we think this is informational severity.

### Alex the Entrepreneurd (judge) commented:

@Sidu28 - if we assumed the Strategy not to be malicious, but to revert for some reason (e.g. lack of liquidity, smart contract bug, etc...)

Would you consider the finding as valid since the user cannot withdrawal from all others due to having one withdraw, or is there some other reason why you believe the finding to be Informational?

### Alex the Entrepreneurd (judge) decreased severity to Medium and commented:

The Warden has shown how, any revert in a strategy (for example it being paused), will make queued withdrawals revert, even if the withdrawal should work for other strategies.

This falls into the category of DOS and I believe is more appropriately judged as Medium.

### ChaoticWalrus (EigenLayer) commented:

@Alex the Entrepreneur - Even if a malicious strategy exists, there is no permanent delay here.

Users can mark the `receiveAsTokens` input to `completeQueuedWithdrawal` as false [here](#) and then the shares will be transferred to the 'withdrawing' user [here](#) who can then *queue a new withdrawal not containing the malicious strategy*.

The description provided by the warden of the `DelegationManager`'s behavior is inaccurate — the funds do not end up 'stuck in the `DelegationManager`' — they will be withdrawable as part of a new queued withdrawal, which excludes the malicious strategy.

[ABA \(warden\) commented:](#)

Hey, <https://github.com/code-423n4/2023-04-eigenlayer/blob/main/src/contracts/core/StrategyManager.sol#L795-L803> does not do:

shares will be transferred to the 'withdrawing' user

It simply increments internal accounting and delegated shares associated to the user in the delegation contract. But these are not withdrawn without a `queueWithdrawal` call.

<https://github.com/code-423n4/2023-04-eigenlayer/blob/main/src/contracts/core/StrategyManager.sol#L647>

*queue a new withdrawal not containing the malicious strategy.*

A new queue can not reuse the already used balance, as it was already accounted for [here](#).

Maybe I am missing something of course. Regardless, awaiting judge feedback.

[ChaoticWalrus \(EigenLayer\) commented:](#)

It simply increments internal accounting and delegated shares associated to the user in the delegation contract. But these are not withdrawn without a

`queueWithdrawal` call

Yes, this is what I would describe as transferring the shares, similar to how the ‘transfer in’ portion of an ERC20 transfer works. But this is a semantic argument; I agree with the substance of your description.

A new queue can not reuse the already used balance as it was already accounted for [here](#)

Maybe I am missing something of course. Regardless, awaiting judge feedback.

I think perhaps you are missing that the internal `_addShares` function calls the `DelegationManager` as well. See [here](#)

If the user marks `receiveAsTokens` as false, then [this line](#) will be triggered, which calls into the internal `_addShares` function and ‘re-adds’ these delegated shares, so they can indeed be queued in a new withdrawal.

I can provide an example if it will help.

[ChaoticWalrus \(EigenLayer\) commented:](#)

Basically, the ‘re-adding’ reverses the accounting taken in the initial `queueWithdrawal` action, to clarify. This then allows the accounting done in the `queueWithdrawal` action to be performed once again when a new withdrawal is queued.

[Alex the Entrepreneur \(judge\) commented:](#)

@ChaoticWalrus - It seems like the worst case scenario would be having to re-queue without the missing strategy, which would require waiting for the withdrawal to unlock.

So the question left to answer is whether an additional wait period is acceptable.

NOTE: I edited this comment because as you pointed out, the shares accounting is internal and doesn’t trigger an interaction with the strategy.

[ChaoticWalrus \(EigenLayer\) commented:](#)



It seems like the worst case scenario would be having to re-queue without the missing strategy which would require waiting for the withdrawal to unlock

Yes, agreed this is the impact + worst-case scenario. The existing functionality was designed, in part at least, to address concerns about malicious Strategies.

So the question left to answer is whether an additional wait period is acceptable

I suppose so, yes. We have deemed this acceptable ourselves, but I could see it being viewed differently. Regardless, the impact here is orders of magnitude less than the original claimed impact.

### [Alex the Entrepreneurd \(judge\) commented:](#)

The Warden has shown how, due to the possibility of queueing a withdrawal with multiple strategies, in the case in which one of the strategies stops working (reverts, paused, malicious), the withdrawal would be denied.

As the sponsor said, in those scenarios, the withdrawer would have to perform a second withdrawal, which would have to be re-queued.

Intuitively, a withdrawer that always withdraws a separate strategy would also never see their withdrawal denied (except for the malicious strategy).

As we can see, there are plenty of side steps to the risk shown. However, the functionality of the function is denied, even if temporarily, leading to it not working as intended, and for this reason I believe Medium Severity to be the most appropriate.

As shown above, the finding could be fixed. However, it seems to me like most users would want to plan around the scenarios described in this finding and its duplicates.



## Low Risk and Non-Critical Issues

For this audit, 15 reports were submitted by wardens detailing low risk and non-critical issues. The [report highlighted below](#) by volodya received the top score from the judge.

The following wardens also submitted reports: [libratus](#), [Cyfrin](#), [niser93](#), [juancito](#), [btk](#), [QiuhaoLi](#), [Aymen0909](#), [Oxnev](#), [ABA](#), [RaymondFam](#), [sashik\\_eth](#), [OxWaitress](#), [ihtishamsudo](#) and [bughunter007](#).



## [L-01] `computePhase0Eth1DataRoot` always returns an incorrect Merkle tree

The Merkle tree creation inside the `computePhase0Eth1DataRoot` function is incorrect.



### Proof of Concept

Provide direct links to all referenced code in GitHub. Add screenshots, logs, or any other relevant proof that illustrates the concept. Not all fields of `eth1DataFields` are being used in an array due to the usage of `i <`

`ETH1_DATA_FIELD_TREE_HEIGHT` instead of `i < NUM_ETH1_DATA_FIELDS`. Check other similar functions.

[src/contracts/libraries/BeaconChainProofs.sol#L160](#)

```
function computePhase0Eth1DataRoot(bytes32[NUM_ETH1_DATA_FIE
    bytes32[] memory paddedEth1DataFields = new bytes32[](2

    for (uint256 i = 0; i < ETH1_DATA_FIELD_TREE_HEIGHT; ++i
        paddedEth1DataFields[i] = eth1DataFields[i];
    }

    return Merkle.merkleizeSha256(paddedEth1DataFields);
}
```



### Recommended Mitigation Steps

```
function computePhase0Eth1DataRoot(bytes32[NUM_ETH1_DATA_FIE
    bytes32[] memory paddedEth1DataFields = new bytes32[](2

    for (uint256 i = 0; i < ETH1_DATA_FIELD_TREE_HEIGHT; ++i
+    for (uint256 i = 0; i < NUM_ETH1_DATA_FIELDS; ++i) {
        paddedEth1DataFields[i] = eth1DataFields[i];
    }
```

```
}  
  
    return Merkle.merkleizeSha256(paddedEth1DataFields);  
}
```



## Assessed type

Math

### Sidu28 (EigenLayer) disagreed with severity and commented:

We believe this is low severity. The code is unused and informally deprecated, but it is indeed technically incorrect.

### Alex the Entrepreneurd (judge) decreased severity to QA and commented:

Agree with the Sponsor, because the code is unused.

### Alex the Entrepreneurd (judge) commented:

Consistently high quality submissions. After grading the QAs I believe the Warden deserves the best place.



**[L-02]** `processInclusionProofKeccak` **does not work as expected**



## Proof of Concept

The function `verifyInclusionKeccak` is not used anywhere but its in the scope of this audit. There is no validation that proof is a tree and a valid tree like it described in the comments. E.x. if proof is less than 32 length, that function will just return a leaf without reverting. In my opinion, function doesn't work as expected and can be exploited. I've submitted the same issue with `processInclusionProofSha256` function that lead to loss a funds for validator due the same issue.

```
function processInclusionProofKeccak(bytes memory proof, byt  
    bytes32 computedHash = leaf;
```

```

    for (uint256 i = 32; i <= proof.length; i+=32) {
        if(index % 2 == 0) {
            // if ith bit of index is 0, then computedHash i
            assembly {
                mstore(0x00, computedHash)
                mstore(0x20, mload(add(proof, i)))
                computedHash := keccak256(0x00, 0x40)
                index := div(index, 2)
            }
        } else {
            // if ith bit of index is 1, then computedHash i
            assembly {
                mstore(0x00, mload(add(proof, i)))
                mstore(0x20, computedHash)
                computedHash := keccak256(0x00, 0x40)
                index := div(index, 2)
            }
        }
    }
    return computedHash;
}

```

## [src/contracts/libraries/Merkle.sol#L49](#)



### Recommended Mitigation Steps

I think its important to add security to that function like this:

```

function processInclusionProofKeccak(bytes memory proof, byt
+     require(proof.length % 32 == 0 && proof.length > 0, "Ir

bytes32 computedHash = leaf;
for (uint256 i = 32; i <= proof.length; i+=32) {
    if(index % 2 == 0) {
        // if ith bit of index is 0, then computedHash i
        assembly {
            mstore(0x00, computedHash)
            mstore(0x20, mload(add(proof, i)))
            computedHash := keccak256(0x00, 0x40)
            index := div(index, 2)
        }
    } else {
        // if ith bit of index is 1, then computedHash i

```

```

        assembly {
            mstore(0x00, mload(add(proof, i)))
            mstore(0x20, computedHash)
            computedHash := keccak256(0x00, 0x40)
            index := div(index, 2)
        }
    }
}
return computedHash;
}

```

[Sidu28 \(EigenLayer\) confirmed](#)

[Alex the Entrepreneur \(judge\) decreased severity to QA and commented:](#)

<https://github.com/code-423n4/2023-04-eigenlayer/blob/398cc428541b91948f717482ec973583c9e76232/src/contracts/operators/MerkleDelegationTerms.sol#L97-L107>

```

// check inclusion of the leafHash in the tree correspor
require(
    Merkle.verifyInclusionKeccak(
        proof,
        merkleRoots[rootIndex].root,
        leafHash,
        nodeIndex
    ),
    "MerkleDelegationTerms.proveEarningsAndWithdraw: pro
);

```

Which calls `processInclusionProofKeccak`

For this reason, I believe the finding to be a Refactoring. Adding the check in the function is a good idea, but the code in scope is safe.

🔗

**[L-03]** `merkleizeSha256` **doesn't work as expected**

<https://github.com/code-423n4/2023-04-eigenlayer/blob/398cc428541b91948f717482ec973583c9e76232/src/contracts/li>



## Proof of Concept

Whenever `merkleizeSha256` is being used in the code, there is always a check that array length is power of 2. E.x.:

```
bytes32[] memory paddedHeaderFields = new bytes32[] (2**BEACON_BI
```

## [contracts/libraries/BeaconChainProofs.sol#L131](#)

But inside the function `merkleizeSha256`, there is no check that incoming array is power of 2.

```
/**
 * @notice this function returns the merkle root of a tree created from
 * @param leaves the leaves of the merkle tree
 *
 * @notice requires the leaves.length is a power of 2
 */
function merkleizeSha256(
    bytes32[] memory leaves
) internal pure returns (bytes32) {
    //there are half as many nodes in the layer above the leaves
    uint256 numNodesInLayer = leaves.length / 2;
    //create a layer to store the internal nodes
    bytes32[] memory layer = new bytes32[] (numNodesInLayer);
    //fill the layer with the pairwise hashes of the leaves
    for (uint i = 0; i < numNodesInLayer; i++) {
        layer[i] = sha256(abi.encodePacked(leaves[2*i], leaves[2*i+1]));
    }
    //the next layer above has half as many nodes
    numNodesInLayer /= 2;
    //while we haven't computed the root
    while (numNodesInLayer != 0) {
        //overwrite the first numNodesInLayer nodes in layer
        for (uint i = 0; i < numNodesInLayer; i++) {
            layer[i] = sha256(abi.encodePacked(layer[2*i], layer[2*i+1]));
        }
        //the next layer above has half as many nodes
        numNodesInLayer /= 2;
    }
}
```

```

    }
    //the first node in the layer is the root
    return layer[0];
}

```

There is a `@notice` that doesn't hold.

**|** `@notice` requires the `leaves.length` is a power of 2

But whenever there is a `require` in `natspec` inside the project, it always holds. E.x.:

```

/**
 * @notice Delegates from `staker` to `operator`.
 * @dev requires that:
 * 1) if `staker` is an EOA, then `signature` is valid ECDSA
 * 2) if `staker` is a contract, then `signature` must will
 */

```

[src/contracts/core/DelegationManager.sol#L89](#)

```

* WARNING: In order to mitigate against inflation/donation
*          minimum amount of shares be either 0 or 1e9. A c
*          be able to withdraw for 1e9-1 or less shares.
*

```

[/src/contracts/strategies/StrategyBase.sol#L72](#)



## Tools Used

You can insert this into remix to check:

```

// SPDX-License-Identifier: GPL-3.0

pragma solidity >=0.7.0 <0.9.0;

import "hardhat/console.sol";

contract Owner {

```

```

mapping(address => bool) internal frozenStatus;
constructor() {
}

function dod() external returns (bytes32){
    bytes32[] memory leaves = new bytes32[](7);
    for (uint256 i = 0; i < 7; ++i) {
        leaves[i] = bytes32(i);
    }
    return merkleizeSha256(leaves);
}

function merkleizeSha256(
    bytes32[] memory leaves
) internal pure returns (bytes32) {
    //there are half as many nodes in the layer above the leaves
    uint256 numNodesInLayer = leaves.length / 2;
    //create a layer to store the internal nodes
    bytes32[] memory layer = new bytes32[](numNodesInLayer);
    //fill the layer with the pairwise hashes of the leaves
    for (uint i = 0; i < numNodesInLayer; i++) {
        layer[i] = sha256(abi.encodePacked(leaves[2*i], leaves[2*i+1]));
    }
    //the next layer above has half as many nodes
    numNodesInLayer /= 2;
    //while we haven't computed the root
    while (numNodesInLayer != 0) {
        //overwrite the first numNodesInLayer nodes in layer
        for (uint i = 0; i < numNodesInLayer; i++) {
            layer[i] = sha256(abi.encodePacked(layer[2*i], layer[2*i+1]));
        }
        //the next layer above has half as many nodes
        numNodesInLayer /= 2;
    }
    //the first node in the layer is the root
    return layer[0];
}
}

```



## Recommended Mitigation Steps

Either remove `@notice` or add this code for more security because sometimes you can just forget to check array size before calling that function:



```

function merkleizeSha256(
    bytes32[] memory leaves
) internal pure returns (bytes32) {
+     uint256 len = leaves.length;
+     while (len > 1 && len % 2 == 0) {
+         len /= 2;
+     }
+     require(len==1, "requires the leaves.length is a power
//there are half as many nodes in the layer above the le
    uint256 numNodesInLayer = leaves.length / 2;
    //create a layer to store the internal nodes
    bytes32[] memory layer = new bytes32[](numNodesInLayer);
    //fill the layer with the pairwise hashes of the leaves
    for (uint i = 0; i < numNodesInLayer; i++) {
        layer[i] = sha256(abi.encodePacked(leaves[2*i], leav
    }
    //the next layer above has half as many nodes
    numNodesInLayer /= 2;
    //while we haven't computed the root
    while (numNodesInLayer != 0) {
        //overwrite the first numNodesInLayer nodes in layer
        for (uint i = 0; i < numNodesInLayer; i++) {
            layer[i] = sha256(abi.encodePacked(layer[2*i], l
        }
        //the next layer above has half as many nodes
        numNodesInLayer /= 2;
    }
    //the first node in the layer is the root
    return layer[0];
}

```

## Remix:

```

// SPDX-License-Identifier: GPL-3.0

pragma solidity >=0.7.0 <0.9.0;

import "hardhat/console.sol";

contract Owner {

    mapping(address => bool) internal frozenStatus;
    constructor() {

```

```

    }

function dod(uint len) external returns (bytes32){
    bytes32[] memory leaves = new bytes32[](len);
    for (uint256 i = 0; i < len; ++i) {
        leaves[i] = bytes32(i);
    }
    return merkleizeSha256(leaves);
}

function merkleizeSha256(
    bytes32[] memory leaves
) internal pure returns (bytes32) {
    uint256 len = leaves.length;
    while (len > 1 && len % 2 == 0) {
        len /= 2;
    }
    require(len==1, "requires the leaves.length is a power of 2");
    //there are half as many nodes in the layer above the leaves
    uint256 numNodesInLayer = leaves.length / 2;
    //create a layer to store the internal nodes
    bytes32[] memory layer = new bytes32[](numNodesInLayer);
    //fill the layer with the pairwise hashes of the leaves
    for (uint i = 0; i < numNodesInLayer; i++) {
        layer[i] = sha256(abi.encodePacked(leaves[2*i], leaves[2*i+1]));
    }
    //the next layer above has half as many nodes
    numNodesInLayer /= 2;
    //while we haven't computed the root
    while (numNodesInLayer != 0) {
        //overwrite the first numNodesInLayer nodes in layer
        for (uint i = 0; i < numNodesInLayer; i++) {
            layer[i] = sha256(abi.encodePacked(layer[2*i], layer[2*i+1]));
        }
        //the next layer above has half as many nodes
        numNodesInLayer /= 2;
    }
    //the first node in the layer is the root
    return layer[0];
}

}

```

[Sidu28 \(EigenLayer\) disputed, disagreed with severity and commented:](#)

The comment is ambiguous, but is intended to actually state a precondition on the input. The comment will be changed.

Alex the Entrepreneurd (judge) decreased severity to QA and commented:

Every instance in the in-scope codebase does check, meaning that the finding cannot be considered a vulnerability.

I can agree with the Warden that a valid refactoring would bring the check in the function to simplify the code.

For this reason, am downgrading to QA - Refactoring (R)



**[L-04]** `claimableUserDelayedWithdrawals` **sometimes returns unclaimable** `DelayedWithdrawals` , **so users will see incorrect data**



## Proof of Concept

The `canClaimDelayedWithdrawal` function will return false for a withdrawal, which the block duration has not passed. The same restriction will be checked whenever an actual withdrawal is triggered, but the `claimableUserDelayedWithdrawals` function does not take into account block duration validation.

```
function claimableUserDelayedWithdrawals(address user) external
    uint256 delayedWithdrawalsCompleted = _userWithdrawals[user]
    uint256 delayedWithdrawalsLength = _userWithdrawals[user]
    uint256 claimableDelayedWithdrawalsLength = delayedWithdrawalsCompleted
    DelayedWithdrawal[] memory claimableDelayedWithdrawals = new DelayedWithdrawal[claimableDelayedWithdrawalsLength]
    for (uint256 i = 0; i < claimableDelayedWithdrawalsLength; i++) {
        claimableDelayedWithdrawals[i] = _userWithdrawals[user][i]
    }
    return claimableDelayedWithdrawals;
```

...

```
function canClaimDelayedWithdrawal(address user, uint256 index) public view returns (bool) {
    return ((index >= _userWithdrawals[user].delayedWithdrawalsLength) ||
        !_userWithdrawals[user].delayedWithdrawals[index].isClaimed);
}
```



## Recommended Mitigation Steps

```
function claimableUserDelayedWithdrawals(address user) external
    uint256 delayedWithdrawalsCompleted = _userWithdrawals[user]
    uint256 delayedWithdrawalsLength = _userWithdrawals[user].length
    uint256 claimableDelayedWithdrawalsLength = delayedWithdrawalsCompleted
    DelayedWithdrawal[] memory claimableDelayedWithdrawals;
    for (uint256 i = 0; i < claimableDelayedWithdrawalsLength; i++) {
        if (block.number < _userWithdrawals[user].delayedWithdrawals[i].blockNumber)
            break;
        claimableDelayedWithdrawals.push(_userWithdrawals[user].delayedWithdrawals[i]);
    }
    return claimableDelayedWithdrawals;
}
```

## [Alex the Entrepreneurd \(judge\) decreased severity to QA](#)



[L-05] The condition for full withdrawals in the code is different from that in the documentation



## Proof of Concept

The condition in [docs](#) for full withdrawal is `validator.withdrawableEpoch < executionPayload.slot/SLOTS_PER_EPOCH` while in the code its `validator.withdrawableEpoch <= executionPayload.slot/SLOTS_PER_EPOCH`.

```
function verifyAndProcessWithdrawal(
    BeaconChainProofs.WithdrawalProofs calldata withdrawalProofs,
    bytes calldata validatorFieldsProof,
    bytes32[] calldata validatorFields,
    bytes32[] calldata withdrawalFields,
    uint256 beaconChainETHStrategyIndex,
    uint64 oracleBlockNumber
) returns (bool) {
    ...
}
```

```
        // reference: uint64 withdrawableEpoch = Endian.fromLittleEndian(validatorFields[BeaconWithdrawalIndex].withdrawableEpoch);
        if (Endian.fromLittleEndianUint64(validatorFields[BeaconWithdrawalIndex].withdrawableEpoch) < withdrawableEpoch) {
            _processFullWithdrawal(withdrawalAmountGwei, validatorIndex);
        } else {
            _processPartialWithdrawal(slot, withdrawalAmountGwei);
        }
    }
}
```

### [src/contracts/pods/EigenPod.sol#L354](#)



#### Recommended Mitigation Steps

Synchronize them with each other.

#### [Sidu28 \(EigenLayer\) disputed, disagreed with severity and commented:](#)

We believe this is an informational-level issue. We failed to update this statement in the higher-level documentation. The code is correct.

#### [Alex the Entrepreneurd \(judge\) decreased severity to QA and commented:](#)

Great catch, but in lack of an impact am downgrading to QA.

Will award extra points. L + 3.



#### [L-06] Missing validation to a threshold value on full withdrawal



#### Proof of Concept

According to the [docs](#) there's supposed to be a validation against a const on full withdrawal, but its missing, which can lead to the system not working as expected.

In this second case, in order to withdraw their balance from the EigenPod, stakers must provide a valid proof of their full withdrawal (differentiated from partial withdrawals through a simple comparison of the amount to a threshold value named `MINFULLWITHDRAWALAMOUNTGWEI`) against a beacon state root.

```

function _processFullWithdrawal(
    uint64 withdrawalAmountGwei,
    uint40 validatorIndex,
    uint256 beaconChainETHStrategyIndex,
    address recipient,
    VALIDATOR_STATUS status
) internal {
    uint256 amountToSend;

    // if the validator has not previously been proven to be
    if (status == VALIDATOR_STATUS.ACTIVE) {
        // if the withdrawal amount is greater than the REQUIRED_BALANCE_GWEI
        if (withdrawalAmountGwei >= REQUIRED_BALANCE_GWEI) {
            // then the excess is immediately withdrawable
            amountToSend = uint256(withdrawalAmountGwei - REQUIRED_BALANCE_GWEI);
            // and the extra execution layer ETH in the context of the podOwner is
            restakedExecutionLayerGwei += REQUIRED_BALANCE_GWEI;
        } else {
            // otherwise, just use the full withdrawal amount
            restakedExecutionLayerGwei += withdrawalAmountGwei;
            // remove and undelegate 'extra' (i.e. "overcommitted") stake
            eigenPodManager.recordOvercommittedBeaconChainETHStrategyIndex(
                validatorIndex,
                beaconChainETHStrategyIndex,
                withdrawalAmountGwei
            );
        }
    }
    // if the validator *has* previously been proven to be 'overcommitted'
    } else if (status == VALIDATOR_STATUS.OVERCOMMITTED) {
        // if the withdrawal amount is greater than the REQUIRED_BALANCE_GWEI
        if (withdrawalAmountGwei >= REQUIRED_BALANCE_GWEI) {
            // then the excess is immediately withdrawable
            amountToSend = uint256(withdrawalAmountGwei - REQUIRED_BALANCE_GWEI);
            // and the extra execution layer ETH in the context of the podOwner is
            restakedExecutionLayerGwei += REQUIRED_BALANCE_GWEI;
            /**
             * since in `verifyOvercommittedStake` the podOwner is
             * in order to allow the podOwner to complete the withdrawal
             */
            eigenPodManager.restakeBeaconChainETH(podOwner,
                amountToSend);
        } else {
            // otherwise, just use the full withdrawal amount
            restakedExecutionLayerGwei += withdrawalAmountGwei;
            /**
             * since in `verifyOvercommittedStake` the podOwner is
             * in order to allow the podOwner to complete the withdrawal
             */
            eigenPodManager.restakeBeaconChainETH(podOwner,
                withdrawalAmountGwei);
        }
    }
}

```

```

// If the validator status is withdrawn, they have already
} else {
    revert("EigenPod.verifyBeaconChainFullWithdrawal: Validator is withdrawn");
}

// set the ETH validator status to withdrawn
validatorStatus[validatorIndex] = VALIDATOR_STATUS.WITHDRAWN;

emit FullWithdrawalRedeemed(validatorIndex, recipient, validatorIndex, amountToSend);

// send ETH to the `recipient`, if applicable
if (amountToSend != 0) {
    _sendETH(recipient, amountToSend);
}
}

```

[src/contracts/pods/EigenPod.sol#L364](#)



## Recommended Mitigation Steps

```

function _processFullWithdrawal(
    uint64 withdrawalAmountGwei,
    uint40 validatorIndex,
    uint256 beaconChainETHStrategyIndex,
    address recipient,
    VALIDATOR_STATUS status
) internal {
+     require(withdrawalAmountGwei >= MIN_FULL_WITHDRAWAL_AMOUNT_GWEI,
+     "stakers must provide a valid proof of their full withdrawal");

    uint256 amountToSend;

    // if the validator has not previously been proven to be active
    if (status == VALIDATOR_STATUS.ACTIVE) {
        // if the withdrawal amount is greater than the required balance
        if (withdrawalAmountGwei >= REQUIRED_BALANCE_GWEI) {
            // then the excess is immediately withdrawable
            amountToSend = uint256(withdrawalAmountGwei - REQUIRED_BALANCE_GWEI);
            // and the extra execution layer ETH in the contract is restaked
            restakedExecutionLayerGwei += REQUIRED_BALANCE_GWEI;
        } else {
            // otherwise, just use the full withdrawal amount
            restakedExecutionLayerGwei += withdrawalAmountGwei;
        }
    }
}

```

```

        // remove and undelegate 'extra' (i.e. "overcomm
        eigenPodManager.recordOvercommittedBeaconChainETH
    }
    // if the validator *has* previously been proven to be '
} else if (status == VALIDATOR_STATUS.OVERCOMMITTED) {
    // if the withdrawal amount is greater than the REQU
    if (withdrawalAmountGwei >= REQUIRED_BALANCE_GWEI) {
        // then the excess is immediately withdrawable
        amountToSend = uint256(withdrawalAmountGwei - RE
        // and the extra execution layer ETH in the cont
        restakedExecutionLayerGwei += REQUIRED_BALANCE_G
        /**
        * since in `verifyOvercommittedStake` the podOv
        * in order to allow the podOwner to complete th
        */
        eigenPodManager.restakeBeaconChainETH(podOwner,
    } else {
        // otherwise, just use the full withdrawal amour
        restakedExecutionLayerGwei += withdrawalAmountGv
        /**
        * since in `verifyOvercommittedStake` the podOv
        * in order to allow the podOwner to complete th
        */
        eigenPodManager.restakeBeaconChainETH(podOwner,
    }
    // If the validator status is withdrawn, they have alrea
} else {
    revert("EigenPod.verifyBeaconChainFullWithdrawal: VF
}

// set the ETH validator status to withdrawn
validatorStatus[validatorIndex] = VALIDATOR_STATUS.WITHI

emit FullWithdrawalRedeemed(validatorIndex, recipient, v

// send ETH to the `recipient`, if applicable
if (amountToSend != 0) {
    _sendETH(recipient, amountToSend);
}
}
}

```

**[Sidu28 \(EigenLayer\) disagreed with severity and commented:](#)**



This is an informational-level issue. We failed to update this statement in the higher-level documentation. This check is not necessary.

[Alex the Entrepreneur \(judge\) decreased severity to QA and commented:](#)

The 4 logical paths seem to cover the possible scenarios.

In lack of further info, am downgrading to QA.



[L-07] User can stake twice on beacon chain from same eipod, thus losing funds due to same withdrawal credentials



## Proof of Concept

There are no restriction to how many times user can stake on beacon with `EigenPodManager` on `EigenPod`, thus all of them will have the same `_podWithdrawalCredentials()` and I think first deposit will be lost.

```
function stake(bytes calldata pubkey, bytes calldata signature) public {
    // stake on ethpos
    require(msg.value == 32 ether, "EigenPod.stake: must init with 32 ether");
    ethPOS.deposit{value : 32 ether}(pubkey, _podWithdrawalCredentials());
    emit EigenPodStaked(pubkey);
}
```

[src/contracts/pods/EigenPod.sol#L159](#)

There are some ways users can make a mistake by calling it twice or they would like to create another one. I've looked into rocketpool contracts; they are not allowing users to stake twice with the same pubkeys, so I think its important to implement the same security issue.

```
function preStake(bytes calldata _validatorPubkey, bytes calldata _signature) public {
    ...
    require(rocketMinipoolManager.getMinipoolByPubkey(_validatorPubkey) == 0, "Minipool already exists");
    // Set minipool pubkey
    rocketMinipoolManager.setMinipoolPubkey(_validatorPubkey, _signature);
}
```

```

// Get withdrawal credentials
bytes memory withdrawalCredentials = rocketMinipoolManager
// Send staking deposit to casper
casperDeposit.deposit{value : prelaunchAmount}(_validatorPubkey, _validatorSignature)
// Emit event
emit MinipoolPrestaked(_validatorPubkey, _validatorSignature);
}

```

## [contracts/contract/minipool/RocketMinipoolDelegate.sol#L235](#)

Same safe thing done in frax finance:

```

function depositEther(uint256 max_deposits) external nonReentrant
...
    // Deposit the ether in the ETH 2.0 deposit contract
    depositContract.deposit{value: DEPOSIT_SIZE}(
        pubkey,
        withdrawalCredential,
        signature,
        depositDataRoot
    );

    // Set the validator as used so it won't get an extra
    activeValidators[pubKey] = true;
...
}

```

## [src/frxETHMinter.sol#L156](#)



### Tools Used

### POC

```

function testWithdrawFromPod() public {
    cheats.startPrank(podOwner);
    eigenPodManager.stake{value: stakeAmount}(pubkey, signature);
+    eigenPodManager.stake{value: stakeAmount}(pubkey, signature);
    cheats.stopPrank();

    IEigenPod pod = eigenPodManager.getPod(podOwner);
    uint256 balance = address(pod).balance;
}

```

```

cheats.deal(address(pod), stakeAmount);

cheats.startPrank(podOwner);
cheats.expectEmit(true, false, false, false);
emit DelayedWithdrawalCreated(podOwner, podOwner, balance);
pod.withdrawBeforeRestaking();
cheats.stopPrank();
require(address(pod).balance == 0, "Pod balance should be 0");
}

```



## Recommended Mitigation Steps

You can look at rocketpool contracts and borrow their logic:

```

function stake(bytes calldata pubkey, bytes calldata signature) public {
+     require(EigenPodManager.getEigenPodByPubkey(_validatorPubkey) == 0, "EigenPod already staked");
+     EigenPodManager.setEigenPodByPubkey(_validatorPubkey, pubkey);

    require(msg.value == 32 ether, "EigenPod.stake: must initiate deposit of 32 ether");
    ethPOS.deposit{value : 32 ether}(pubkey, _podWithdrawalContract);
    emit EigenPodStaked(pubkey);
}

```

### [ChaoticWalrus \(EigenLayer\) disagreed with severity and commented:](#)

We believe this is purely informational. People can already stake multiple times through the `ETH2Deposit` contract.

### [Alex the Entrepreneur \(judge\) decreased severity to QA and commented:](#)

Downgrading to QA for now. @volodya - please do send me proof that a new withdrawal would be bricked (I believe it would be releasable via the normal flow).

### [Alex the Entrepreneur \(judge\) commented:](#)

Have had a informal confirmation that excess ETH is refunded as rewards.

In lack of additional info, am maintaining the judgment.



# Gas Optimizations

For this audit, 13 reports were submitted by wardens detailing gas optimizations. The [report highlighted below](#) by neutiyo0 received the top score from the judge.

*The following wardens also submitted reports: [OxSmartContract](#), [turvy\\_fuzz](#), [pontifex](#), [niser93](#), [Oxnev](#), [QiuhaoLi](#), [Aymen0909](#), [ihtishamsudo](#), [tonisives](#), [clayj](#), [ReyAdmirado](#) and [naman1778](#).*



## [G-01] Optimize `merkleizeSha256` function for gas-efficiency

Although the current implementation of the [merkleizeSha256](#) function in the [Merkle](#) contract is correct, it can be more gas-efficient by making use of the following optimizations:



### 1. In-place Computation

The `merkleizeSha256` function can be optimized by using in-place computation to store intermediate hashes at each level of the Merkle tree. This approach eliminates the need to create new arrays, reducing memory usage and gas costs.

Note: this optimization requires the `leaves` array not to be used again after it is modified. Based on the current implementation, this optimization is safe because the `leaves` array is not used again after it is modified.



### 2. Assembly

The use of assembly code to load the left and right siblings into memory is more gas-efficient than using the `abi.encodePacked` function.



### 3. Unchecked Arithmetic

The use of unchecked arithmetic for `uint i` is more gas-efficient as it skips checks for overflow or underflow. This optimization is safe because `i` is always less than `numNodesInLayer`, meaning that overflow is not possible.



### Proof of Concept

The function `merkleizeSha256Optimized` provided below is an optimized version of the `merkleizeSha256` function.

### [src/contracts/libraries/MerkleOptimized.sol](#)

```
// SPDX-License-Identifier: MIT
pragma solidity =0.8.12;

library MerkleOptimized {
    /**
     * @notice Returns the Merkle root of a tree created from
     * @param leaves The leaves of the Merkle tree. This par
     * @return The Merkle root of the tree.
     *
     * @notice Requires the leaves.length is a power of 2.
96  * @dev This is adapted from https:
97  */
98  function merkleizeSha256Optimized(
99      bytes32[] memory leaves
100 ) internal pure returns (bytes32) {
101     // Reserve memory space for our hashes.
102     bytes memory buf = new bytes(64);
103
104     // We'll need to keep track of left and right sibling
105     bytes32 leftSibling;
106     bytes32 rightSibling;
107
108     // Number of non-empty nodes at the current depth.
109     uint256 rowSize = leaves.length;
110
111     // Common sub-expressions
112     uint256 halfRowSize; // rowSize / 2
113
114     while (rowSize > 1) {
115         halfRowSize = rowSize / 2;
116
117         for (uint256 i = 0; i < halfRowSize; ) {
118             leftSibling = leaves[(2 * i)];
119             rightSibling = leaves[(2 * i) + 1];
120             assembly {
121                 mstore(add(buf, 32), leftSibling)
122                 mstore(add(buf, 64), rightSibling)
123             }
124
```

```

125         leaves[i] = sha256(buf);
126
127         unchecked {
128             ++i;
129         }
130     }
131
132     rowSize = halfRowSize;
133 }
134
135 return leaves[0];
136 }
137 }

```

| Function Name            | min  | avg    | median | max     | # calls |  |
|--------------------------|------|--------|--------|---------|---------|--|
| merkleizeSha256          | 2353 | 274987 | 62975  | 1396167 | 10      |  |
| merkleizeSha256Optimized | 2136 | 238896 | 56158  | 1197190 | 10      |  |

| Improvement |        |  |
|-------------|--------|--|
| Minimum     | 9.22%  |  |
| Average     | 13.12% |  |
| Median      | 10.82% |  |
| Maximum     | 14.25% |  |

The data shows a significant increase in gas efficiency with the use of `merkleizeSha256Optimized` compared to `merkleizeSha256`. It's worth emphasizing, that these results are influenced by the input data and execution environment, so the actual improvement may differ in other contexts. Nonetheless, the results provide valuable insight into the potential gas cost savings that can be achieved by leveraging the optimized version of the function.

The test codes are the following:

`src/test/unit/Merkle.t.sol`

```

// SPDX-License-Identifier: BUSL-1.1
pragma solidity =0.8.12;

```

```

import {Test} from "forge-std/Test.sol";
import {Merkle} from "../..//contracts/libraries/Merkle.sol";
import {MerkleOptimized} from "../..//contracts/libraries/MerkleC

contract MerkleMock {
    function merkleizeSha256(
        bytes32[] calldata leaves
    ) external pure returns (bytes32) {
        return Merkle.merkleizeSha256(leaves);
    }

    function merkleizeSha256Optimized(
        bytes32[] calldata leaves
    ) external pure returns (bytes32) {
        return MerkleOptimized.merkleizeSha256Optimized(leaves);
    }
}

contract MerkleTest is Test {
    MerkleMock public c;

    function setUp() external {
        c = new MerkleMock();
    }

    function gen(uint256 length) internal pure returns (bytes32[]
        bytes32[] memory leaves = new bytes32[](length);
        for (uint i = 0; i < length; i++) {
            leaves[i] = bytes32(i);
        }
        return leaves;
    }

    function testMerkleizeSha256Equivalence() external {
        for (uint i = 2; i <= 1024; i *= 2) {
            assertEq(
                c.merkleizeSha256(gen(i)),
                c.merkleizeSha256Optimized(gen(i)),
                "ok"
            );
        }
    }
}

```



## Recommendation

Consider optimizing `merkleizeSha256` by using in-place computation, assembly, and unchecked arithmetic.



## [G-02] Use unchecked arithmetic in `processInclusionProofKeccak` and `processInclusionProofSha256` functions

The `processInclusionProofKeccak` and `processInclusionProofSha256` functions in the `Merkle` contract include unnecessary arithmetic checks for incrementing `uint256 i` in a for-loop. By using unchecked arithmetic, the gas cost of executing these functions can be reduced.



### 1. `processInclusionProofKeccak`

[src/contracts/libraries/Merkle.sol#L48-L50](#)

```
function processInclusionProofKeccak(bytes memory proof, bytes32 computedHash = leaf;
    for (uint256 i = 32; i <= proof.length; i+=32) {
    ...
```



### 2. `processInclusionProofSha256`

[src/contracts/libraries/Merkle.sol#L99-L101](#)

```
function processInclusionProofSha256(bytes memory proof, bytes32[1] memory computedHash = [leaf];
    for (uint256 i = 32; i <= proof.length; i+=32) {
    ...
```

Based on the current implementation, this optimization is safe because overflow is not possible, as the length of `proof` is validated before the function call.



## Recommendation



Consider using unchecked arithmetic for `uint256 i`.

```
unchecked {  
    i += 32;  
}
```



## Disclosures

C4 is an open organization governed by participants in the community.

C4 Audits incentivize the discovery of exploits, vulnerabilities, and bugs in smart contracts. Security researchers are rewarded at an increasing rate for finding higher-risk issues. Audit submissions are judged by a knowledgeable security researcher and solidity developer and disclosed to sponsoring developers. C4 does not conduct formal verification regarding the provided code but instead provides final verification.

C4 does not provide any guarantee or warranty regarding the security of this project. All smart contract software should be used at the sole risk and responsibility of users.

Top

An open organization | [Twitter](#) | [Discord](#) | [GitHub](#) | [Medium](#) | [Newsletter](#) | [Media kit](#) | [Careers](#) | [code4rena.eth](#)